

Running MITgcm forward and adjoint on TACC machines and sverdrup

An T. Nguyen, Patrick Heimbach, Helen Pillar, Kylie Kinne, Kirstin Schulz, Matthew Goldberg

January 25, 2023

0 Modules, optfiles, iddev:

0.1 Modules

Prior to discussing setting up the specific domains, we need to check that the modules appropriate for the MITgcm are loaded. This is my list of modules on [stampede2](#), which is loaded by default in my [.bashrc](#), which is available at [/home1/03901/atnguyen/.bashrc_stampede2](#).

```
login1.stampede2 $ module list
```

Currently Loaded Modules:

```
1) autotools/1.1      2) cmake/3.16.1      3) xalt/2.10.2
4) TACC                5) intel/18.0.2      6) libfabric/1.7.
7) impi/18.0.2        8) python2/2.7.15    9) netcdf/4.3.3.1   10) git/2.24.1
```

On [Frontera](#), these are my modules in [/home1/03901/atnguyen/.bashrc_frontera](#):

```
cd ~/
```

```
-rwx----- 1 atnguyen G-816140 2379 Jan 22 2020 .bashrc_frontera
```

```
login3.frontera(1004)$ module list
```

Currently Loaded Modules:

```
1) autotools/1.2      3) pmix/3.1.4         5) xalt/2.10.2      7) intel/19.1.1     9) python3/3.7.0    11) git/2.24.1
2) cmake/3.16.1      4) hwloc/1.11.12     6) TACC              8) impi/19.0.9     10) netcdf/4.6.2
```

0.2 Optfiles

In addition we have also created optfiles that link appropriate libraries and compilers for the skylake nodes on stampede2 and frontera:

```
/work/03901/atnguyen/BitBucket/computing/optfiles/
```

```
linux_amd64_ifort+mpi_stampede2_aste
```

```
linux_amd64_ifort+mpi_frontera_aste1
```

On sverdrup the optfiles are at:

```
/home/atnguyen/nansen/computing/optfiles/
```

```
linux_amd64_ifort+mpi_sverdrup_seaice_tides
```

```
linux_amd64_idort+mpi_sverdrup_aste
```

```
optfile_mvapich2_sverdrup
```

37 **0.3 Interactive node**

38 To ask for an interactive skylake node on stampede2 for 2 hours (max), the command is:

39 `idev -A atn-startup -p skx-dev -t 2:0:0`

40 To run mpi process on the interactive node, the syntax is:

41 `ibrun -n ${nprocs} ./mitgcmuv${forwadj}`

Model Domain Configurations

There are five domain configurations covered in this note: (1) regional **ASTE1080**, (2) **global1080**, (3) regional **ASTE4320**, (4) **ASTE270**, and (5) **ASTE90** (miniASTE). The number “1080” in (1) and (2) refers to a grid we use that has 4×1080 points along the earth’s equator (covering 360° longitudes at $\sim 110\text{km}/1^\circ$). This translates to $\sim 360^\circ / (4 \times 1080) \times (110\text{km}/1^\circ) = 9\text{km}$ horizontal grid spacing at the equator. For ASTE4320, ASTE270 and ASTE90, the horizontal grid spacings correspond to 5.3km, 36.7km and 110km

1 ASTE1080

1.1 Domain and domain decomposition

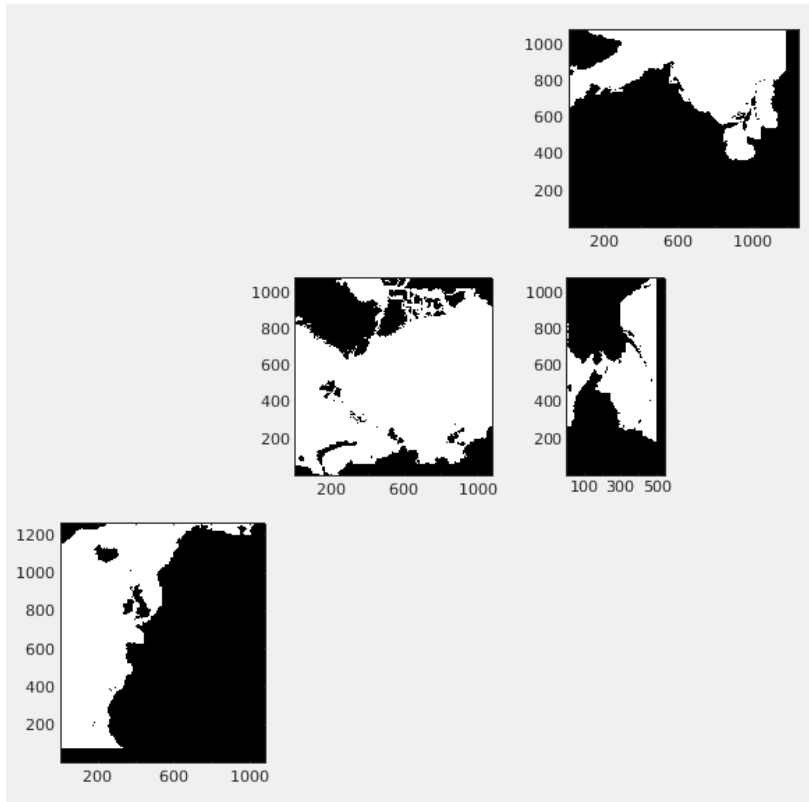


Figure 1: ASTE1080 domain

This is a regional domain, containing a total of 1080×4140 horizontal grid points and 90 vertical levels, yielding a total of 402,408,000 grid points. **Fig. 1** shows the full ASTE1080 domain. Note that the domain is **not** a rectangle of size 1080×4140 but rather is composed of several “faces” (four in **Fig. 1**) such that when re-assembled back into a rectangle the size is 1080×4140 . The decomposition of the domain is done along the horizontal slice (non-continuous 1080×4140 grid points, see **Fig. 1**) via tile partitioning. An example of the decomposition into tiles of size 90×90 is shown in **Fig. 2**. Once decomposition is finished, the MITgcm allows for exclusion of “land” tiles, such that we only solve for “wet tiles”, e.g., less than the maximum number of $1080 \times 4140 / 90 = 552$ tiles and less than the total 402,408,000 grid points.




Figure 2: ASTE1080 decomposed into 90x90 tiles. The gray area is “blanked out” such that only the white area, a total 320 tiles (shown as incremental numbers, instead of 552 tiles), will be solved for. For this decomposition, the total number of grid points will be $90 \times 90 \times 320 \times 90$ -vertical levels = 233,280,000 grid points, which is $\sim 57\%$ of the total number of grid points for the entire domain. The smaller the tile size, e.g., 18x18 or 15x15, the more gray areas we can remove from the computation. With this 90x90 decomposition, a total of 320 tiles translates to 320 CPUs is needed, and that each CPU will solve for $90 \times 90 = 8100$ grid points in the horizontal (plus grid overlaps, and also to be multiplied with number of vertical levels). This yields a max of 8100×90 -vertical levels = 729,000 grid points per CPU, which can be taxing for the CPU memory if the state is stored. Memory demands can be reduced by using smaller tile size, with the trade-off being the increased communications between CPUs and nodes.

Table 1 shows some of the possible decompositions and associated total number of CPUs required. For a forward calculation, tile sizes $\sim 30 \times 30$ or larger are OK. For the adjoint calculation, due to the memory required to store the state, tile sizes 15×15 , 18×18 , or 20×20 are needed.

sNx	sNy	No. of tiles = No. of CPUs (nProc)
90	90	320
90	60	463
60	60	659
45	60	860
45	45	1149
sqrt(512 45)	36	1430
36	36	1784
36	30	2099
30	30	2484
20	30	3671
20	20	5438
20	18	6027
18	18	6695
15	15	9405

Table 1: Decomposition of the ASTE1080 regional domain into various tile sizes. $sN[x,y]$ are the number of grid points in the $[x,y]$ -dir per tile. Note that the tiles need not be squared, but likely for optimal communication/exchange with edges of different tile, i.e. different CPUs, it's desirable to keep the tiles close to a square. There are more possibilities than what's listed here, but these are some samples for testing. Horizontal lines are only drawn to improve ease of reading.

1.2 Location of files

The main directory for this configuration is at:

```

/work/03901/atnguyen/MITgcm_c67/MITgcm/
/work/03901/atnguyen/MITgcm_c67/mysetups/aste_1080x1260x540x90/
  code_tides_mds_prof_sp1/
  build_tides_mds_prof_sp1/
  input_tides/
  aste1080_tides.bash
/work/03901/atnguyen/BitBucket/computing/optfiles/
  linux_amd64_ifort+mpi_stampede2_aste
  linux_amd64_ifort+mpi_frontera_aste1

```

To compile in your own directory, copy all of the above to the appropriate location. Below, I use an example with **\$LOCALDIR** being **/scratch/03901/atnguyen/localdir**:

```

export atnwork='/work/03901/atnguyen/'
export LOCALDIR='/scratch/03901/atnguyen/localdir/'
cd $LOCALDIR

```

```

77 cp -rp $atnwork/MITgcm_c67/MITgcm/ ./MITgcm
78 cp -rp $atnwork/MITgcm_c67/mysetups/aste_1080x1260x540x90/code_tides_mds_prof_sp1 ./code
79 cp -rp $atnwork/MITgcm_c67/mysetups/aste_1080x1260x540x90/input_tides ./NAMELISTS
80 cp -p $atnwork/MITgcm_c67/mysetups/aste_1080x1260x540x90/aste1080_tides.bash ./aste1080.bash
81 cp -p $atnwork/BitBucket/computing/optfiles/linux_amd64_ifort+mpi_stampede2_aste ./
82 cp -p $atnwork/BitBucket/computing/optfiles/linux_amd64_ifort+mpi_frontera_aste1 ./
83 cp -p /home1/03901/atnguyen/.bashrc_stampede2 ./bashrc_stampede2
84 S

```

1.3 Instruction to compile and run the model

1.3.1 Compiling:

The compilation of the code requires the main-branch **MITgcm** directory and the specific **code** for the ASTE1080 domain, as well as the appropriate optfile depending on whether we're using stamped2 or frontera,

```

90 linux_amd64_ifort+mpi_stampede2_aste
91 linux_amd64_ifort+mpi_frontera_aste1

```

Before compiling, choose the tile decomposition by reading off **Table. 1** and pick the row you want, for example, **90x90x320**. The first manual step is to cp the **SIZE.h** file associated with this tile set in **code/**:

```

94 cd $LOCALDIR/code
95 cp SIZE_h_90x90x320 SIZE.h

```

Next, create the **build** directory and compile the MITgcm:

```

97 cd $LOCALDIR
98 mkdir build
99 cd build
100 ../MITgcm/tools/genmake2 -mpi -of=../linux_amd64_ifort+mpi_stampede2_aste -mods=../code -rd=../MITgcm
101 make depend
102 make -j 4

```

where the **"genmake2"** command creates the file **Makefile** in the directory which allows you to check all the module lists and compare to mine, and the last "make" command creates a **"mitgcmuv"** executable and the **"-j 4"** will compile in parallel to speed things up, using 4 CPUs in this example. One last step I do is to move this executable to match the chosen tile decomposition:

```

107 mv mitgcmuv mitgcmuv_90x90x320

```

NOTE: if you have to **re-compile**, in the same **build** dir, there are two levels of clearing you need:

a) If all modules are correct, and you're only changing **SIZE.h** to recompile, then the file **Makefile** can be re-used as follows:

```

111 cd $LOCALDIR/build
112 make CLEAN
113 make makefile (note small case 'm' in 'makefile')
114 make depend
115 make -j 4

```

b) If your modules are incorrect, you cannot use the existing **Makefile** that was created in the **build/** dir and will simply need to empty the dir and recompile from scratch using **genmake2**.

118 1.3.2 Prepare and submit jobscript:

119 Next, edit the jobscript file `aste1080.bash`. Line 7 to choose the account:

```
120 7 #SBATCH -A atn-startup
```

121 Next, un-comment the correct option associated with the choice `90x90x320` in this example:

```
122 10 #SBATCH -N 7
```

```
123 11 #SBATCH -n 336
```

```
124 59 nprocs=320
```

```
125 60 snx=90
```

```
126 61 sny=90
```

127 Finally, edit the line to point to where your `$LOCALDIR` is:

```
128 89 scratchdir=/scratch/03901/atnguyen/localdir
```

129 Now submit the job:

```
130 sbatch aste1080.bash
```

131 I've a successfully created all of the above steps and have an example run using `90x90x320` at:

```
132 /scratch/03901/atnguyen/localdir/
```

```
133 linux_amd64_ifort+mpi_frontera_aste1
```

```
134 linux_amd64_ifort+mpi_stampede2_aste
```

```
135 MITgcm/
```

```
136 code/
```

```
137 build/
```

```
138 NAMELISTS/
```

```
139 aste1080.bash
```

```
140 run_hourly_pk0001577880/ <-- the succesful run
```

141 1.4 Expected results

142 In the run directory, you should be able to see progress and monitored physical oceanographic statistics and
143 any warnings in:

```
144 /scratch/03901/atnguyen/localdir/run_hourly_pk0001577880/
```

```
145 STDOUT.0000, STDERR.0000
```

146 The run is set to output at quite high frequency, with the output files dumped in the subdirectory

```
147 /scratch/03901/atnguyen/localdir/run_hourly_pk0001577880/; diags/
```

2 Global1080

2.1 Domain and domain decomposition

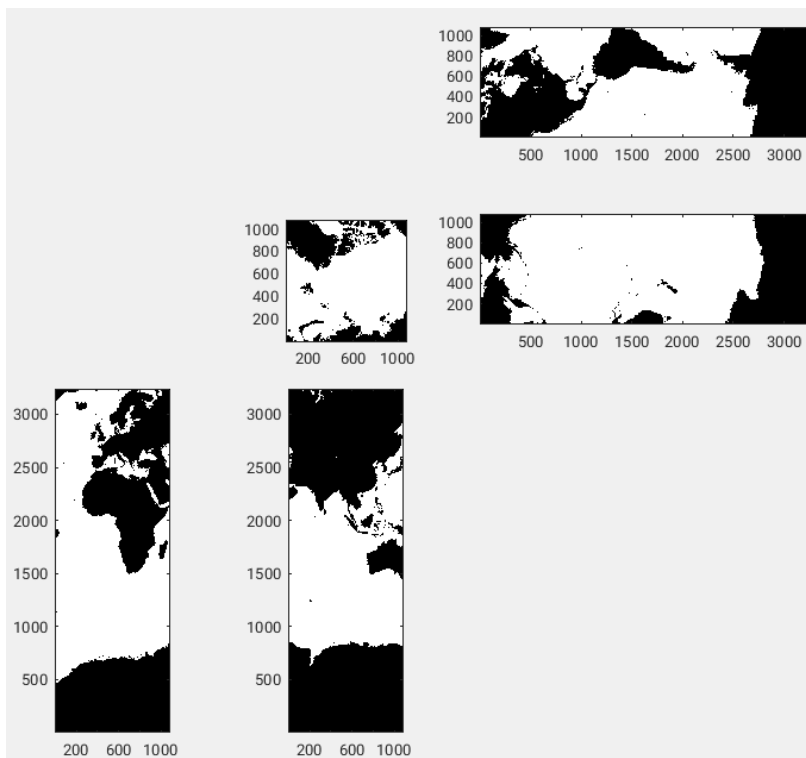


Figure 3: global1080 domain

This is a global domain, containing a total of $1080 \times (1080 \times 13)$ horizontal grid points and 90 vertical levels, yielding a total of $15,163,200 \times 90 = 1,364,688,000$ grid points. **Fig. 3** shows the full global1080 domain. Note that the domain is composed of five “faces” of sizes $1080 \times (1080 \times 3)$, $1080 \times (1080 \times 3)$, 1080×1080 , $(1080 \times 3) \times 1080$, $(1080 \times 3) \times 1080$, such that when re-assembled back into a rectangle the size is $1080 \times (1080 \times 13)$. The decomposition of this domain is similar to that for the **ASTE1080**, and **Table. 2** shows the number of CPUs required for various sizes of tiles for **global1080**.

The number of nodes and syntax for requesting resources (last 3 columns of **Table. 2**) are calculated as follows. Assume X is the number of cpus per node ($X=48$ on stampede2, $X=56$ on frontera), we get, for the example of $120 \times 120 \times 757$:

on stampede2:

$$757/X = 15.77$$

$16 \times X = 768$ = total number of nodes (16) and cpus (768) requested

$768 - 757 = 11$ = offset in the first node to allow for more memory on cpu 0

on frontera:

$$757/X = 13.52$$

$14 \times X = 784$ = total number of nodes (14) and cpus (784) requested

$784 - 757 = 27$ = offset in the first node to allow for more memory on cpu 0

sNx	sNy	No. of tiles = No. of CPUs (nProc)	Frontera(minimum)			Frontera at 48CPU/node		
			#node	#cpus asked	offset	#node	#cpus asked	offset
120	120	757	14	784	27	16	896	11
90	90	1,300	24	1344	44	28	1568	44
90	72	1,597	29	1624	27	34	1904	35
90	60	1,900	35	1925	25	40	2240	20
72	72	1,979	36	2016	37	42	2352	37
72	60	2,349	43	2365	27	49	2744	3
60	60	2,787	50	2800	21	59	3304	45
60	45	3,667	66	3696	29	77	4312	29
45	45	4,831	87	4872	41	101	5656	17
45	36	5,982	107	5992	10	125	7000	18
36	36	7,418	133	7448	30	155	8680	22
36	30	8,837	158	8848	11	185	10360	43
30	30	10,540	189	10584	44	220	12320	20
20	30	15,625	280	15680	55	326	18256	23
20	20	23,135	414	23184	49	482	26992	1
20	18	25,631	458	25648	17	534	29904	1
18	18	28,407	508	28448	41	592	33152	9
18	15	33,967	607	33992	25	708	39648	17
15	15	40,570	725	40600	30	846	47376	38
12	15	50,477	902	50512	35	1052	58912	19
12	12	62,818	1122	62832	14	1309	73304	14
12	9	83,321	1488	83328	7	1736	97216	7
9	9	110,551	1975	110600	49	2304	129024	41
9	6	164,943	2946	164976	33	3437	192472	33
6	6	246,058	4394	246064	6	5127	287112	38

Table 2: Decomposition of the **global1080** regional domain into various tile sizes. sN[x,y] are the number of grid points in the [x,y]-dir per tile. It is desirable to keep the tiles close to a square. **NOTE: Toward the bottom of this Table, as the No. of tiles increases, in `code/W2_EXCH2_SIZE.h`, one needs to ensure the variable `W2_maxNbTiles` \geq Number of CPUs required. If not, the model will crash during run-time with the error reporting this problem in `run/STDERR.0000`.** For Frontera, $\#node = \text{round_up}(\#CPUs / 56)$, $\#cpus \text{ asked} = \#node * 56$, and $offset = \#cpus \text{ asked} \text{ minus } \#cpus \text{ required}$. Syntax for jobscript will be: `#SBATCH -N #node; #SBATCH -n #cpus asked; ibrun -n #cpus -o offset mitgcmuv` . Thus for the example of **120x120x757**, the syntax is `#SBATCH -N 14; #SBATCH -n 784; ibrun -n 757 -o 27 mitgcmuv` .

A= CPUs/node desired	N=ceil(nProc/A) (No. of nodes requested)	o=N*A-nProc Offset in the 1st node	No. CPUs requested = N*56	Frontera syntax
56	14	27	784	-N 14 -n 784 -o 27
52	15	23	840	-N 15 -n 840 -o 23
48	16	11	896	-N 16 -n 896 -o 11
44	18	35	1008	-N 18 -n 1008 -o 35
40	19	3	1064	-N 19 -n 1064 -o 3
36	22	35	1232	-N 22 -n 1232 -o 35
32	24	11	1344	-N 24 -n 1344 -o 11
28	28	27	1568	-N 28 -n 1568 -o 27
24	32	11	1792	-N 32 -n 1792 -o 11

Table 3: Breakdown of what to ask for -N, -n, and -o for **frontera** with 56 cpus / node, assuming we're using the tile decomposition **sNx by sNy = 120x120** and requiring **nProc=757** (first row in Table. 2).

2.2 Location of files

For the **global1080**, running can be either on frontera or stamped2. As both machines can see **/work/** , the input binaries for the configuration are stored there. For that reason, we need to copy them, about 170GB total, to **/scratch** (on stamped2) or **/scratch1** (on frontera) prior to copying the **code** and **NAMELISTS**:

```
export atnwork='/work/03901/atnguyen/'
export LOCALDIR='/scratch/03901/atnguyen/global_llc1080/'
cd $LOCALDIR
mkdir jra55 jra55_do
cp -p $atnwork/jra55/*_2010 ./jra55
cp -p $atnwork/jra55/*_2011 ./jra55
cp -p $atnwork/jra55_do/*_2010 ./jra55_do
cp -p $atnwork/jra55_do/*_2011 ./jra55_do
cp -rp $atnwork/llc1080/global/run_template ./
```

Next we proceed with copying the **code**, **NAMELISTS**, optfiles, jobscript:

```
cd $LOCALDIR
cp -rp $atnwork/MITgcm_c67/MITgcm/ ./MITgcm
cp -rp $atnwork/MITgcm_c67/mysetups/global_llc1080/code_tides_mds_prof_sp1 ./code
cp -rp $atnwork/MITgcm_c67/mysetups/global_llc1080/input_tides ./NAMELISTS
cp -p $atnwork/MITgcm_c67/mysetups/global/glob1080/*.bash ./
cp -p $atnwork/BitBucket/computing/optfiles/linux_amd64_ifort+mpi_stamped2_aste ./
cp -p $atnwork/BitBucket/computing/optfiles/linux_amd64_ifort+mpi_frontera_aste1 ./

cp -p /home1/03901/atnguyen/.bashrc_stamped2 ./bashrc_stamped2
OR
cp -p /home1/03901/atnguyen/.bashrc_frontera ./bashrc_frontera
```

2.2.1 Compiling:

The compilation of the code requires the main-branch **MITgcm** directory and the specific **code** for the glob1080 domain, as well as the appropriate optfile depending on whether we're using stamped2 or frontera,

```

196     linux_amd64_ifort+mpi_stampede2_aste
197     linux_amd64_ifort+mpi_frontera_aste1

```

Before compiling, we need to choose the tile decomposition. Read off **Table. 2** and pick the row you want, for example, **120x120x757**. The first manual step is to cp the **SIZE.h** file associated with this tile set in **code/**:

```

201     cd $LOCALDIR/code
202     cp SIZE_h_120x120x757 SIZE.h

```

Next, check to make sure in **code/W2_EXCH2.SIZE.h** the variable **W2_maxNbTiles** \geq Number of CPUs required.

```

205     cd $LOCALDIR/code
206     grep -n 'W2_maxNbTiles = ' W2_EXCH2_SIZE.h
207
208     35:C      W2_maxNbTiles = Nb of active tiles (=nSx*nSy*nPx*nPy) + Max_Nb_BlankTiles
209     43:      PARAMETER ( W2_maxNbTiles = nSx*nSy*nPx*nPy * 2 + 1200 )

```

For **120x120x757**, we must ensure **W2_maxNbTiles = nSx*nSy*nPx*nPy * 2 + 1200 \geq 757**. If not, the model will crash during run-time with the error reporting this problem in **run/STDERR.0000**.

Next, create the **build** directory and compile the MITgcm (note choose correct optfile depending on the machine, stampede2 in this example):

```

215     cd $LOCALDIR
216     mkdir build
217     cd build
218     ../MITgcm/tools/genmake2 -mpi -of=../linux_amd64_ifort+mpi_stampede2_aste -mods=../code -rd=../MITgcm
219     make depend
220     make -j 4

```

where the “**genmake2**” command creates the file **Makefile** in the **build/** directory which allows you to check all the module lists and compare to mine, and the last “make” command creates a “**mitgcmuv**” executable and the “-j 4” will compile in parallel to speed things up, using 4 CPUs in this example. One last step I do is to move this executable to match the chosen tile decomposition:

```

225     mv mitgcmuv mitgcmuv_120x120x757

```

NOTE: if you have to **re-compile**, in the same **build** dir, there are two levels of clearing you need:

a) If all modules are correct, and you’re only changing **SIZE.h** to recompile, then the file **Makefile** can be re-used as follows:

```

229     cd $LOCALDIR/build
230     make CLEAN
231     make makefile (note small case ‘m’ in ‘makefile’)
232     make depend
233     make -j 4

```

2.2.2 Prepare and submit jobscript:

Next, edit the jobscript file `glob1080_stampede2.bash` (or `glob1080_fr1.bash` if running on frontera).
Line 7 to choose the account:

```
7 #SBATCH -A atn-startup
```

Choose the node type, **skx-normal** for stampede2, and **normal** for frontera, and how long the job will run:

```
6 #SBATCH -t 8:00:00
```

Next, un-comment the correct option associated with the choice **120x120x757** in this example:

```
15 #SBATCH -N 16
```

```
16 #SBATCH -n 768
```

```
243
```

```
35 nprocs=757
```

```
36 snx=120
```

```
37 sny=120
```

Finally, edit the line to point to where your `$LOCALDIR` is:

```
51 scratchdir=/scratch/03901/atnguyen/global_llc1080
```

Now submit the job (example for stampede2):

```
sbatch glob1080_stampede2.bash
```

I've successfully created all of the above steps and finished a short (4-hour wall time) example run using **120x120x757** at:

```
/scratch/03901/atnguyen/global_llc1080/  
  jra55/  
  jra55_do/  
  run_template/  
  linux_amd64_ifort+mpi_frontera_aste1  
  linux_amd64_ifort+mpi_stampede2_aste  
  MITgcm/  
  code/  
  build/  
  NAMELISTS/  
  glob1080_stampede2.bash  
  glob1080_fr1.bash  
  run_hourly_pk0001577880/
```

3 ASTE4320

3.1 Domain and decomposition

To resolve the high frequency processes in the Arctic, we need a configuration based on the llc4320 grid (e.g., $\sim 1\text{km}$ in the Arctic). The **ASTE4320** domain is configured to include the entire North Atlantic Ocean above $\sim 5^\circ\text{N}$ (**Fig. 4**). The decomposition of this domain is shown in **Table. 4**. This is a configuration that **has never been tested** due to lack of resources, both the computation as well as storage. Specifically, to spin up a configuration such as this would require access to enough nodes at efficient time to start from scratch with minimum time-step, the ramp up to a stable configuration. A typical start-up and spin-up will require at least 10-15 runs to address all instabilities/issues.

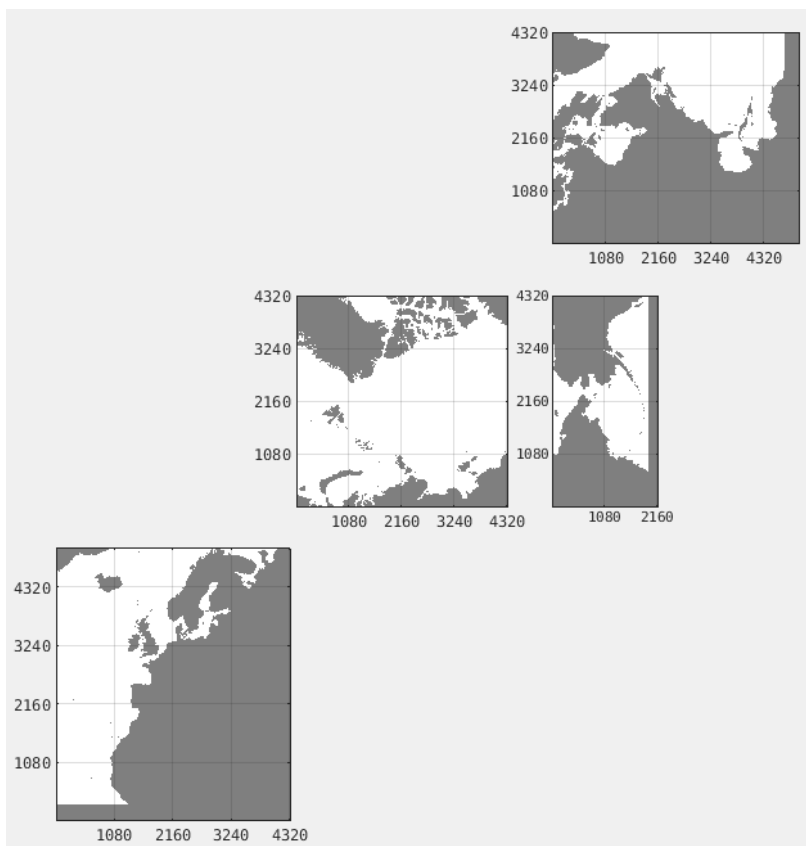


Figure 4: ASTE4320 domain

sNx	sNy	No. of tiles = No. of CPUs (nProc)	Frontera		
			#node	#cpus asked	offset
120	120	2812	51	2856	44
90	90	4810	86	4816	6
90	60	7075	127	7112	37
60	60	10428	187	10472	44
45	60	13746	246	13776	30
45	45	18139	324	18144	5
36	36	27874	498	27888	14
36	30	33251	594	33264	13
30	30	39660	709	39704	44
30	24	49345	882	49392	47
24	24	61426	1097	61432	6
20	24	73276	1309	73304	28
20	20	87541	1564	87584	43
18	20	97044	1733	97048	4
18	18	107651	1923	107688	37
18	15	128655	2298	128688	33
15	15	153825	2747	153832	7
12	15	191600	3422	191632	32
12	12	238624	4262	238672	48

Table 4: Decomposition of the **ASTE4320** regional domain into various tile sizes. $sN[x,y]$ are the number of grid points in the $[x,y]$ -dir per tile. It is desirable to keep the tiles close to a square. **NOTE: Toward the bottom of this Table, as the No. of tiles increases, in [code/W2_EXCH2_SIZE.h](#), one needs to ensure the variable [W2_maxNbTiles](#) \geq Number of CPUs required. If not, the model will crash during run-time with the error reporting this problem in [run/STDERR.0000](#).** For Frontera, $\#node = \text{round_up}(\#CPUs / 56)$, $\#cpus \text{ asked} = \#node * 56$, and $offset = \#cpus \text{ asked} \text{ minus } \#cpus \text{ required}$. Syntax for jobscript will be: `#SBATCH -N #node; #SBATCH -n #cpus asked; ibrun -n #cpus -o offset mitgcmuv .` Thus for the example of **120x120x2812**, the syntax is `#SBATCH -N 51; #SBATCH -n 2856; ibrun -n 2812 -o 44 mitgcmuv .`

4 ASTE270

4.1 Domain and decomposition

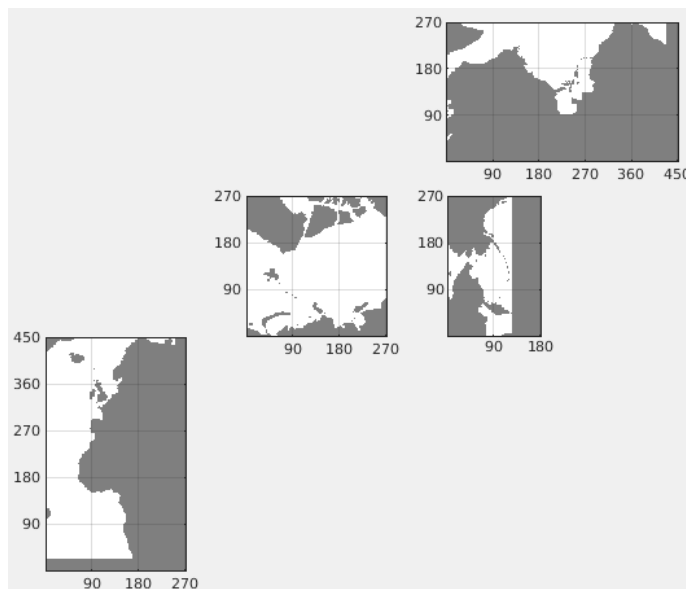


Figure 5: ASTE270 domain

sNx	sNy	No. of tiles = No. of CPUs (nProc)
90	90	36
45	45	117
30	30	242
18	18	580

Table 5: Decomposition of the ASTE270 regional domain into various tile sizes. $sN[x,y]$ are the number of grid points in the $[x,y]$ -dir per tile.

4.2 Location of files

```

$ /work/03901/atnguyen/BitBucket/computing/optfiles/linux_amd64_ifort+mpi_stampede2_aste
$ /work/03901/atnguyen/MITgcm_c65q/MITgcm/

export astedir='/work/03901/atnguyen/MITgcm_c65q/mysetups/aste_270x450x180/'
$ astedir/run_RERUN_ASTE_R1_lfs_adv30_mombudg10day/
  code_ASTE_it62/
  code_adv7_bypass_tamc/
  code_HP_mombudg/
  NAMELISTS/
  script_ASTE_R1_mombudg.bash

```

```

288
289     /scratch/projects/ecco/
290     llc270/aste_270x450x180/run_template/
291     forcing/jra55

```

292 4.2.1 Compiling:

```

293     export LOCALDIR='/scratch/03901/atnguyen/aste_270x450x180/mombudg/'
294     mkdir $LOCALDIR
295     cd $LOCALDIR
296     cp -p $/work/03901/atnguyen/BitBucket/computing/optfiles/linux_amd64_ifort+mpi_stampede2_aste ./
297     cp -rp /work/03901/atnguyen/MITgcm_c65q/MITgcm ./
298     cp -rp $astedir/code_ASTE_it62 ./
299     cp -rp $astedir/code_adv7_bypass_tamc ./
300     cp -rp $astedir/code_HP_mombudg ./

```

301 Refer to **Table. 5** to choose tile size. Here we use **90x90x36** as an example:

```

302     cd $LOCALDIR
303     cd code_ASTE_it62
304     cp SIZE_h_90x90x36_ol4 SIZE.h
305
306     cd $LOCALDIR
307     mkdir build
308     cd build
309
310     ../MITgcm/tools/genmake2 -of=../linux_amd64_ifort+mpi_stampede2_aste \
311     '-mods=../code_ASTE_it62 ../code_adv7_bypass_tamc ../code_HP_mombudg' -rd=../MITgcm
312     make depend
313     make -j 4
314
315     mv mitgcmuv mitgcmuv_90x90x36

```

316 4.2.2 Prepare and submit jobscript:

```

317     cd $LOCALDIR
318     cp -rp $astedir/NAMELISTS ./
319     cp -p /work/03901/atnguyen/MITgcm_c65q/mysetups/aste_270x450x180/script_ASTE_R1_mombudg.bash ./

```

320 Edit **NAMELISTS/data** to comment out advection scheme 7 and choose advection scheme 30 (make
321 sure the hash sign # is at the **first** column):

```

322     38 ## code_adv7_bypass_tamc and set overlap ol[x,y]=8 in SIZE.h
323     39 # tempAdvScheme=7,
324     40 # saltAdvScheme=7,
325     41 #
326     42 ## if use adv30, only need overlap ol[x,y]=4
327     43 tempAdvScheme=30,
328     44 saltAdvScheme=30,

```


329 Edit [NAMELISTS/data.pkg](#) to turn off ecco and profiles to reduce computational cost as follows:

```
330      10 useECCO      = .FALSE.,
331      12 useProfiles  = .FALSE.,
```

332 Edit jobscript [script_ASTE_R1_mombudg.bash](#). For the choice **90x90x36** the 36 CPUs fit within one
333 node, so we can use the skx-dev node to test:

```
334      6 #SBATCH -p skx-dev
335      7 #SBATCH -t 1:00:00
336
337     13 #SBATCH -N 1
338     14 #SBATCH -n 48
339
340     47 #SBATCH -A atn-startup
341
342     63 nprocs=36
343     64 snx=90
344     65 sny=90
345
346    103 criosdir=/scratch/projects/ecco/llc270/aste_270x450x180/
347    104 scratchdir=/scratch/03901/atnguyen/aste_270x450x180/mombudg/
```

348 Now submit the job:

```
349      sbatch script_ASTE_R1_mombudg.bash
```

350 I've a successfully created all of the above steps and finished a very short example run using **90x90x36** at:

```
351      /scratch/03901/atnguyen/aste_270x450x180/mombudg/
352      linux_amd64_ifort+mpi_stampede2_aste
353      MITgcm/
354      code_ASTE_it12/
355      code_adv7_bypass_tamc/
356      code_HP_mombudg/
357      build/
358      NAMELISTS/
359      script_ASTE_R1_mombudg.bash
360      run_ASTE_R1_mombudg_pk0000000007/  <-- mombudg outputs in diags/
```

5 ASTE90

5.1 Domain and decomposition

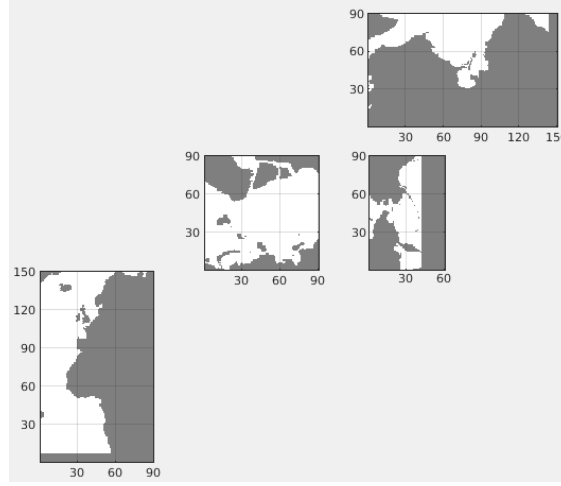


Figure 6: ASTE90 domain

sNx	sNy	No. of tiles = No. of CPUs (nProc)
30	30	36
15	30	64
15	15	117

Table 6: Decomposition of the ASTE90 regional domain into various tile sizes. $sN[x,y]$ are the number of grid points in the $[x,y]$ -dir per tile.

5.2 Location of files

```

/work/03901/atnguyen/BitBucket/computing/optfiles/linux_amd64_ifort+mpi_stampede2_aste
/work/03901/atnguyen/MITgcm_c65q/MITgcm/

/work/03901/atnguyen/MITgcm_c65q/mysetups/aste_90x150x60/
  code_horflux_boxcont_FW5_logdiffkr_mod2_diffkh_fixedlogical_adxOFF
  code_HP_mombudg/
  input_boxcont_fw_beaufort_adxOFF_mombudg
  script_aste90_mombudg.bash

/scratch/projects/ecco/
  llc90_nontelelescope/aste_90x150x90/run_template/
  forcing/jra55
  forcing/jra55_it12xx

/scratch/03901/atnguyen/aste_90x150x60/mombudg/

```

5.2.1 Compiling:

Note: Sometimes copy-and-pasting these lines results in an error due to the mistranslation of the apostrophe. Just in case, it is good to write each line out by hand.

```
export astedir='/work/03901/atnguyen/MITgcm_c65q/mysetups/aste_90x150x60/'
export LOCALDIR='/scratch/03901/atnguyen/aste_90x150x60/mombudg/'
mkdir $LOCALDIR
cd $LOCALDIR
cp -p /work/03901/atnguyen/BitBucket/computing/optfiles/linux_amd64_ifort+mpi_stampede2_aste ./
cp -rp /work/03901/atnguyen/MITgcm_c65q/MITgcm ./
cp -rp $astedir/code_horflux_boxcont_FW5_logdiffkr_mod2_diffkh_fixedlogical_adxOFF ./code
cp -rp $astedir/code_HP_mombudg ./

cd $LOCALDIR/code
cp SIZE_h_30x30x36 SIZE.h

cd $LOCALDIR
mkdir build
cd build

../MITgcm/tools/genmake2 -of=../linux_amd64_ifort+mpi_stampede2_aste \
'-mods=../code ../code_HP_mombudg' -rd=../MITgcm
make depend
make -j 4
mv mitgcmuv mitgcmuv_30x30x36
```

5.2.2 Prepare and submit jobscript:

```
cd $LOCALDIR
cp -rp $astedir/input_boxcont_fw_beaufort_adxOFF_mombudg ./NAMELISTS
cp -p $astedir/script_aste90_mombudg.bash ./
```

Edit [NAMELISTS/data.pkg](#) to make sure pkg ecco, ctrl, profiles, autodiff, and smooth are off to reduce computational cost as follows:

```
10 useECCO      = .FALSE.,
11 useCTRL      = .FALSE.,
12 useProfiles  = .FALSE.,
13 useAUTODIFF  = .FALSE.,
14 useSMOOTH    = .FALSE.,
```

Edit jobscript [script_aste90_mombudg.bash](#). For the choice **30x30x36** the 36 CPUs fit within one node, so we can use the skx-dev node to test:

```
6 #SBATCH -p skx-dev
7 #SBATCH -t 1:00:00
9 #SBATCH -N 1
10 #SBATCH -n 36
```

```

419     14 #SBATCH -A atn-startup
420     30 nprocs=36
421     31 snx=30
422     32 sny=30
423     50 criosdir=/scratch/projects/ecco/llc90_nontele Scop/aste_90x150x60/
424     51 scratchdir=/scratch/03901/atnguyen/aste_90x150x60/mombudg/

```

425 Now submit the job:

```

426     sbatch script_aste90_mombudg.bash

```

427 I've a successfully created all of the above steps and finished a very short example run using **30x30x36** at:


```

428     /scratch/03901/atnguyen/aste_90x150x60/mombudg/
429     linux_amd64_ifort+mpi_stampede2_aste
430     MITgcm/
431     code/
432     code_HP_mombudg/
433     build/
434     NAMELISTS/
435     script_aste90_mombudg.bash
436     run_aste90_mombudg_10days/ <-- mombudg outputs in diags/

```

437

438



439

440

449

450

461

463

464

```

466  ../MITgcm/tools/genmake2 -mpi -of=../linux_amd64_ifort+mpi_stampede2_aste \
467      '-mods=../code ../code_HP_mombudg' -rd=../MITgcm
468  make depend
469  make -j 4
470
471  Building using knl node:
472  ../MITgcm/tools/genmake2 -mpi -of=../linux_amd64_ifort+mpi_stampede2_aste_knl \
473      '-mods=../code ../code_HP_mombudg' -rd=../MITgcm
474  make depend
475  make -j 4
476
477  Edit the jobscript script\_labsea\_adxOFF\_mombudg.bash, where the development queue for KNL is
478  “development” and for SKX “skx-dev”. Below is the example for KNL node:
479
480  6 #SBATCH -p development
481  9 #SBATCH -t 1:00:00
482  10 #SBATCH -N 1
483  11 #SBATCH -n 1
484  12 #SBATCH -A atn-startup
485
486  44 scratchdir=/scratch/03901/atnguyen/labsea/mombudg
487  34 node="knl"
488  47 builddir=$scratchdir/build_${node}
489  49 workdir=$scratchdir/run_mombudg_${node}
490
491  Save and submit the job:
492
493  sbatch script_labsea_adxOFF_mombudg.bash
494
495  There are two runs provided, only 10-time-steps, using knl and skx:
496
497  /scratch/03901/atnguyen/labsea/mombudg/
498  run_mombudg_knl
499  run_mombudg_skx
500
501  A comparison of time can be seen to compare the time difference between knl and skx:
502
503  login1.stampede2(2085)$ more run_mombudg_knl/run.MITGCM.timing
504  Mon Jun  7 01:09:17 CDT 2021
505  Mon Jun  7 01:10:01 CDT 2021 <-- 44sec
506
507  login1.stampede2(2086)$ more run_mombudg_skx/run.MITGCM.timing
508  Mon Jun  7 01:07:34 CDT 2021
509  Mon Jun  7 01:07:50 CDT 2021 <-- 16sec!
510
511  44/16=2.75
512
513  The longer we run, the potentially slower the run using knl nodes can become. However, knl nodes are
514  readily available , unlike skx which is very limited.

```

505 6.2.1 Expected results:

```
506 /scratch/03901/atnguyen/labsea/mombudg/  
507 run_mombudg_knl/diags/  
508 run_mombudg_skx/diags/
```

509 6.2.2 Some tacc tips

510 How to keep a variable between login sessions: The best way is to create a shortcut in `/.bashrc_stampede2`

```
511 export $LOCALHOST='/scratch/07117/ni2kita/labsea/mombudg/'  
512 save it, then source  
513  
514 source ~/.bashrc
```

515 Note about running on scratch and what to copy to work:

```
516 cd $LOCALDIR  
517 tar czvf build_skx.tgz build_skx/ <-- "c" create, "z" zip, "v" verbose, "f" file  
518 tar czvf build_knl.tgz build_knl/  
519  
520 mkdir $workdir/labsea/  
521 mkdir $workdir/labsea/mombudg  
522 cp build_skx.tgz $workdir/labsea/mombudg/  
523 cp build_knl.tgz $workdir/labsea/mombudg/  
524  
525 cd $workdir/labsea/mombudg/  
526 tar xzvf build_skx.tgz build_skx/mitgcmuv  
527 tar xzvf build_skx.tgz build_skx/Makefile  
528  
529 tar xzvf build_knl.tgz build_knl/mitgcmuv  
530 tar xzvf build_knl.tgz build_knl/Makefile
```